

# COMP9415 Review

---

## 2D Transformations

---

- **Affine transformations**

- **Translation**

- Translation is the process of moving an object in space

- $$\begin{pmatrix} q_1 \\ q_2 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & \phi_1 \\ 0 & 1 & \phi_2 \\ 0 & 0 & 1 \end{pmatrix}$$

- **Rotation**

- Rotate objects around the origin

- $$\begin{pmatrix} q_1 \\ q_2 \\ 1 \end{pmatrix} = \begin{pmatrix} \cos(\theta) & -\sin(\theta) & 0 \\ \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

- **Scaling**

- Scale along both axes.

- $$\begin{pmatrix} s_x p_1 \\ s_y p_2 \\ 1 \end{pmatrix} = \begin{pmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

- **Shear**

- Shear is the unwanted child of affine transformations.

- Horizontal

- $$\begin{pmatrix} q_1 \\ q_2 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & h & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} p_1 \\ p_2 \\ 1 \end{pmatrix}$$

- Vertical

- $$\begin{pmatrix} q_1 \\ q_2 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ v & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} p_1 \\ p_2 \\ 1 \end{pmatrix}$$

- **Transformation pipeline**

- $$P_{camera} \xleftarrow{view} P_{world} \xleftarrow{model} P_{local}$$

- The model transform transforms points in the local coordinate system to the world coordinate system

- The view transform transforms points in the world coordinate system to the camera's coordinate system


- **Matrix**

- $P_{world} = Trans(Rot(Scale(P_{camera})))$
- The view matrix:  $P_{camera} = Scale^{-1}(Rot^{-1}(Trans^{-1}(P_{camera})))$

- **Decomposing**

- $Translation = (\phi_1, \phi_2, 1)^T$
- $Rotation = atan2(i_2, i_1)$
- $Scale = |i|$
- **Perpendicular** (垂直):  $i \cdot j = 0$

$$\begin{pmatrix} i_1 & j_1 & \phi_1 \\ i_2 & j_2 & \phi_2 \\ 0 & 0 & 1 \end{pmatrix}$$



axes                      origin

- **Camera**

- In addition to the transformation properties inherited from SceneObject the Camera has an **aspect ratio**.

## 3D Transformations

---

- **TransTranslation**

- $M_T = \begin{pmatrix} 1 & 0 & 0 & \phi_1 \\ 0 & 1 & 0 & \phi_2 \\ 0 & 0 & 1 & \phi_3 \\ 0 & 0 & 0 & 1 \end{pmatrix}$

- **Scale**

- $M_T = \begin{pmatrix} S_x & 0 & 0 & 0 \\ 0 & S_y & 0 & 0 \\ 0 & 0 & S_z & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$

- **Shear**

- $M_T = \begin{pmatrix} 1 & h & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$

- **Rotate**

- Right Hand Rule: For any axis, if the right thumb points in the positive direction of the axis the right fingers curl in the direction of rotation

- **Rotate X**

- $M_x = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos(\theta) & -\sin(\theta) & 0 \\ 0 & \sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$

- **Rotate Y**

- $M_y = \begin{pmatrix} \cos(\theta) & 0 & \sin(\theta) & 0 \\ 0 & 1 & 0 & 0 \\ -\sin(\theta) & 0 & \cos(\theta) & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$

- **Rotate Z**

- $M_z = \begin{pmatrix} \cos(\theta) & -\sin(\theta) & 0 & 0 \\ \sin(\theta) & \cos(\theta) & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$

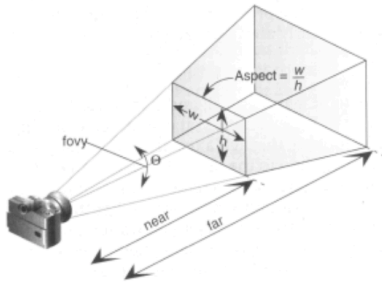
- **Projection**

- Projection happens after the model and view transformations have been applied, so all points are in camera coordinates.
- Points with negative z values in camera coordinates are in front of the camera.
- **Canonical View Volume (CVV)**

- **Perspective**

- $M_{perspective} = \begin{pmatrix} n & 0 & 0 & 0 \\ 0 & n & 0 & 0 \\ 0 & 0 & a & b \\ 0 & 0 & -1 & 0 \end{pmatrix}$

- This matrix is **not affine**.
-

| Situation   | Matrix  |
|---|---|
|  | $M_P = \begin{pmatrix} \frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & \frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ 0 & 0 & -\frac{(f+n)}{f-n} & -\frac{2fn}{f-n} \\ 0 & 0 & -1 & 0 \end{pmatrix}$ |

## Draw Algorithm

- Painter's algorithm
  - Sort geometric primitives by depth
  - Draw in order from back to front
  - **Problem**
    - Intersect polygon
- BSP Tree
  - One possible solution is to use Binary Space Partitioning trees (BSP trees)
  - They recursively divide the world into polygons that are behind or in front of other polygons and split polygons when necessary.
  - Then it is easy to traverse and draw polygons in a front to back order
  - Building the tree is slow and it needs to be rebuilt every time the geometry changes.
  - Best for rendering static geometry where tree can just be loaded in.
- Depth Buffer
  - Initially the depth buffer is initialised to 1 (maximum depth in window coords).
  - Each polygon is drawn fragment by fragment.
  - If it is closer, we update the pixel in the colour buffer and update the buffer value to the new pseudodepth. If not we discard it.

## Triangles

|  |  |  |  |    |    |    |    |    |    |
|--|--|--|--|----|----|----|----|----|----|
|  |  |  |  |    |    |    |    |    |    |
|  |  |  |  |    |    |    |    |    |    |
|  |  |  |  |    |    |    |    |    |    |
|  |  |  |  |    |    |    |    |    |    |
|  |  |  |  |    |    |    |    |    |    |
|  |  |  |  | .7 | .7 | .7 | .7 | .7 | .7 |
|  |  |  |  | .6 | .6 | .6 | .6 | .6 | .6 |
|  |  |  |  | .5 | .5 | .5 | .5 | .5 | .5 |
|  |  |  |  | .4 | .4 | .4 | .4 | .4 | .4 |
|  |  |  |  | .3 | .3 | .3 | .3 | .3 | .3 |
|  |  |  |  | .2 | .2 | .2 | .2 | .2 | .2 |

## Buffer

|    |    |    |    |    |    |    |    |    |    |
|----|----|----|----|----|----|----|----|----|----|
| .1 | 1  | 1  | 1  | 1  | 1  | 1  | 1  | 1  | 1  |
| .1 | .1 | 1  | 1  | 1  | 1  | 1  | 1  | 1  | 1  |
| .2 | .2 | .2 | 1  | 1  | 1  | 1  | 1  | 1  | 1  |
| .2 | .2 | .2 | .2 | 1  | 1  | 1  | 1  | 1  | 1  |
| .3 | .3 | .3 | .3 | .3 | .7 | .7 | .7 | .7 | .7 |
| .3 | .3 | .3 | .3 | .3 | .3 | .6 | .6 | .6 | .6 |
| .4 | .4 | .4 | .4 | .4 | .4 | .4 | .5 | .5 | .5 |
| .4 | .4 | .4 | .4 | .4 | .4 | .4 | .4 | .4 | .4 |
| .5 | .5 | .5 | .5 | .5 | .3 | .3 | .3 | .3 | .3 |
| .5 | .5 | .5 | .5 | .2 | .2 | .2 | .2 | .2 | .2 |

## Limitations

- **Ambient light**

- $I_{ambient} = I_a \rho_a$ 
  - $I_a$  is the ambient light intensity
  - $\rho_a$  is the ambient reflection coefficient  
in the range (0,1) (usually  $\rho_a = \rho_d$ )
- Lighting with just diffuse and specular lights gives very stark shadows.
- In reality shadows are not completely black.
- It is too computationally expensive to model this in detail.
- Light is coming from all directions, reflected off other objects, not just from 'sources'

- **Diffuse light (Lambert's Cosine Law)**

- $I_d = I_s \rho_d (\hat{s} \cdot \hat{m})$ 
  - $I_s$  is the source intensity
  - $\rho_d$  is the diffuse reflection coefficient in [0,1]
- Both vectors are normalised
- When the angle is 0 degrees the cosine is 1
  - All the reflected light back
- When the angle is 90 degrees
  - None of the light is reflected back
- When the angle is > 90 degrees
  - $\cos$  gives us a negative value! This is not what we want.

- **Specular light (Phong model)**

- $\hat{r} = -\hat{s} + 2(\hat{s} \cdot \hat{m})\hat{m}$

- $I_{sp} = \max(0, I_s \rho_{sp} (\hat{r} \cdot \hat{v})^f)$ 
  - $\rho_{sp}$  is the specular reflection coefficient in the range [0,1]
  - $f$  is the phong exponent, typically in the range [1,128]
  - Larger values of the Phong exponent  $f$  make  $\cos(\Phi)^f$  smaller, produce less scattering, creating more mirror-like surfaces.
- **Blinn Phong Model**
  - $I_{sp} = \max(0, I_s \rho_{sp} (\hat{h} \cdot \hat{m})^f)$
  - halfway vector  $\hat{h} = \frac{\hat{s} + \hat{v}}{2}$
  - Phong/Blinn Phong model only reflects light sources, not the environment.
  - It is good for adding bright highlights but cannot create a true mirror.
- **Total intensity for the vertex**
  - $I = I_{ambient} + I_d + I_{sp}$
  - $I = I_a \rho_a + \max(0, I_s \rho_d (\hat{r} \cdot \hat{v})) + \max(0, I_s \rho_{sp} (\hat{r} \cdot \hat{v})^f)$
- **Spotlights**
  - A spotlight has a direction and a cutoff angle
  - Spotlights are also attenuated, so the brightness falls off as you move away from the centre.
  - $I = I_s (\cos(\beta))^\varepsilon$ 
    - $\varepsilon$  is the attenuation factor

## Shading

---

- **Flat shading**
  - Calculated for each face
  - The simplest option is to shade the entire face the same colour
  - **Pro**
    - Diffuse illumination
    - For flat surfaces with distant light sources
    - Non-realistic/retro rendering
    - The fastest shading option
  - **Con**
    - close light sources
    - specular shading
    - curved surfaces
- **Gouraud shading**
  - Calculated for each vertex and interpolated for every fragment
  - Illumination is calculated at each of these vertices.
  - Gouraud shading is a simple smooth shading model.
  - We calculate fragment colours by bilinear interpolation on neighbouring vertices.
  - **Pro**

- curved surfaces
  - close light sources
  - diffuse shading
- **con**
  - more expensive than flat shading
  - handles specular highlights poorly
- **Phong shading**
  - Calculated for every fragment
  - designed to handle specular lighting better than Gouraud.
  - It also handles diffuse better as well.
  - illumination values are calculated per fragment rather than per vertex
  - **Pro**
    - Handles specular lighting well
    - Improves diffuse shading
    - More physically accurate
  - **Con**
    - Slower than Gouraud as normals and illumination values have to be calculated per pixel rather than per vertex.
    - In the old days this was a BIG issue. Not so much any more.

|         | Where    | Pro                                | Con                                  |
|---------|----------|------------------------------------|--------------------------------------|
| Flat    | Face     | Flat surfaces, a retro blocky look | Curved surfaces, specular highlights |
| Gouraud | Vertex   | Curved surfaces, diffuse shading   | Specular highlights                  |
| Phong   | Fragment | Diffuse and specular shading       | Old hardware                         |

## Color

- **Alpha blending**

- $$\begin{pmatrix} r \\ g \\ b \end{pmatrix} \leftarrow a \begin{pmatrix} r_{image} \\ g_{image} \\ b_{image} \end{pmatrix} + (1 - a) \begin{pmatrix} r \\ g \\ b \end{pmatrix}$$

- **Example:**

- If the pixel on the screen is currently green, and we draw over it with a red pixel, with alpha = 0.25

$$\begin{aligned} \blacksquare p &= \begin{pmatrix} 0 \\ 1 \\ 0 \\ 1 \end{pmatrix} p_{image} = \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0.25 \end{pmatrix} \\ \blacksquare p &= 0.25 \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0.25 \end{pmatrix} + 0.75 \begin{pmatrix} 0 \\ 1 \\ 0 \\ 1 \end{pmatrix} = \begin{pmatrix} 0.25 \\ 0.75 \\ 0 \\ 0.8125 \end{pmatrix} \end{aligned}$$

- **Con**

- Alpha blending depends on the order that pixels are drawn.
- You need to draw transparent polygons after the polygons behind them.
- Must be Back-to-front order

## Curves

---

- **de Casteljau Algorithm** (Bézier curves )

- Linear interpolation

- $P(t) = (1 - t)P_0 + tP_1$

- Quadratic interpolation

- $P(t) = (1 - t)^2 P_0 + 2t(1 - t)P_1 + t^2 P_2$

- Cubic interpolation

- $P(t) = (1 - t)^3 P_0 + 3t(1 - t)^2 P_1 + 3t^2(1 - t)P_2 + t^3 P_3$

- **Bézier curves**

- $\sum_{k=0}^m B_k^m(t)P_k$

- $m$  is the degree of the curve

- $P_0 \dots P_m$  are the control points

- The coefficient functions are called Bernstein polynomials.

- $B_k^m(t) = \binom{m}{k} t^k (1 - t)^{m-k}$

- $\binom{m}{k} = \frac{m!}{k!(m-k)!}$  is binomial function

- **example**

- $m = 3$

- $B_0^3(t) = (1 - t)^3$

- $B_1^3(t) = 3t(1 - t)^2$

- $B_2^3(t) = 3t^2(1 - t)$

- $B_3^3(t) = t^3$

- $P(t) = (1 - t)^3 P_0 + 3t(1 - t)^2 P_1 + 3t^2(1 - t)P_2 + t^3 P_3$

- **Tangents**

- Tangent vector

- $\frac{dP(t)}{dt} = \sum_{k=0}^m \frac{dB_k^m(t)}{dt} P_k$



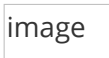
- $\frac{dP(t)}{dt} = \sum_{k=0}^{m-1} \frac{dB_k^{m-1}(t)}{dt} (P_{k+1} - P_k)$
- **L-System** (Lindenmayer System)
  - Can give us realistic plants and trees
  - L-system is a formal grammar
  - Symbols
    - A, B, +, -
    - $A \rightarrow B - A - B$
    - $B \rightarrow A + B + A$
  - Use a **LIFO** stack to save and restore global state like position and heading

## Textures

---

- Usage
  - 1. Load or creating textures
  - 2. Passing the texture to a shader
  - 3. Mapping texture co-ordinates to vertices
- Texture WRAP
  - GL.GL\_REPEAT (default)
  - GL.GL\_MIRRORED\_REPEAT
  - GL.GL\_CLAMP\_TO\_EDGE

```
gl.glTexParameteri( GL.GL_TEXTURE_2D, GL.GL_TEXTURE_WRAP_S,
GL.GL_REPEAT);
gl.glTexParameteri( GL.GL_TEXTURE_2D, GL.GL_TEXTURE_WRAP_T,
GL.GL_REPEAT);
```

- **Textures and shading**
  - The simplest approach is to replace illumination calculations with a texture look-up.
    - $I(P) = T(s(P), t(P))$
    - This produces objects which are not affected by lights or color.
  - A more common solution is to use the texture to modulate the ambient and diffuse reflection coefficients.
    - $I(P) = T(s, t)[I_a \rho_a + I_d \rho_d(\hat{s} \cdot \hat{m})] + I_s \rho_s(\hat{r} \cdot \hat{v})^f$
    - We usually leave the specular term unaffected because it is unusual for the material colour to affect specular reflections.
- **Magnification** (zoom in)
  - Nearest Texel
    - Find the nearest texel
    - 
  - Bilinear Filtering
    - Find the nearest four texels and use bilinear interpolation over them
- **Minification** (zoom out)

- **Aliasing**
  - It occurs when samples are taken from an image at a lower resolution than repeating detail in the image.
- **Filtering**
  - One screen pixel overlaps multiple texels but is taking its value from only one of those texels.
  - A better approach is to average the texels that contribute to that pixel.
  - Doing this on the fly is expensive.
- **MIP mapping**
  - Starting with a 512x512 texture we compute and store 256x256, 128x128, 64x64, 32x32, 16x16, 8x8, 4x4, 2x2 and 1x1 versions.
  - This takes total memory = 4/3 original.

```
gl.glGenerateMipmap(GL.GL_TEXTURE_2D);
```

  - The simplest approach is to use the next smallest mipmap for the required resolution.
- **Trilinear filtering**
  - A more costly approach is trilinear filtering
    - Use bilinear filtering to compute pixel values based on the next highest and the next lowest mipmap resolutions.
    - Interpolate between these values depending on the desired resolution.
- **Aniso Filtering**
  - If a polygon is on an oblique angle away from the camera, then minification may occur much more strongly in one dimension than the other.
  - Anisotropic filtering is filtering which treats the two axes independently.
- **RIP Mapping**
  - RIP mapping is an extension of MIP mapping which down-samples each axis and is a better approach to anisotropic filtering
    - 256x256 image has copies at: 256x128, 256x64, 256x32, 256x16, ..., 128x256, 128x128, 128x64, ..., 64x256, 64x128, etc.
  - **Con**
    - Does not handle diagonal anisotropy.
    - More memory required for RIP maps (4 times as much).
    - Not implemented in OpenGL
- **Multi-texturing**
  - Have to pass two different textures to the shader.
  - two different sets of texture coordinates
- **Animated textures**
  - Animated textures can be achieved by loading multiple textures and using a different one on each frame.
- **Rendering to a texture**
  - A common trick is to set up a camera in a scene, render the scene into an offscreen

buffer, then copy the image into a texture to use as part of another scene.

- **Reflection mapping (cube mapping)**

- Doing this in general is expensive, but we can make a reasonable approximation with textures
  - Generate a cube that encloses the reflective object.
  - Place a camera at the centre of the cube and render the outside world onto the faces of the cube.
  - Use this image to texture the object
- To apply the reflection-mapped texture to the object we need to calculate appropriate texture coordinates.
- We do this by tracing a ray from the camera, reflecting it off the object and then calculating where it intersects the cube.
- **Pros**
  - Produces reasonably convincing polished metal surfaces and mirrors
- **Cons**
  - Expensive: Requires 6 additional render passes per object
  - Angles to near objects are wrong.
  - Does not handle self-reflections or recursive reflections.

- **Shadow buffering**

- keep a shadow buffer for each light source.
- The shadow buffer is like the depth buffer, it records the distance from the light source to the closest object in each direction.
  - Render the scene from each light's viewpoint capturing only z-info in shadow (depth) buffer (color buffer turned off)
  - Render the scene from camera's point of view, using the previously captured shadow buffers to modulate the fragments
- When rendering a point P
  - Project the point into the light's clip space.
  - Calculate the index (i,j) for P in the shadow buffer
  - Calculate the pseudodepth d relative to the light source
  - If  $\text{shadow}[i,j] < d$  then P is in the shadow
- **Pro**
  - Provides realistic shadows
  - No knowledge or processing of the scene geometry is required
- **Cons**
  - More computation
  - Shadow quality is limited by precision of shadow buffer. This may cause some aliasing artefacts.
  - Shadow edges are hard.
  - The scene geometry must be rendered once per light in order to generate the shadow map for a spotlight, and more times for an omnidirectional point light.

- **Light Mapping**

- If our light sources and large portions of the geometry are static then we can precompute the lighting equations and store the results in textures called light maps.
- This process is known as baked lighting.
- **Pro**
  - Sophisticated lighting effects can be computed at compile time, where speed is less of an issue.
- **Con**
  - Memory and loading times for many individual light maps.
  - Not suitable for dynamic lights or moving objects.
  - Potential aliasing effects depending on the resolution of the light maps.
- **Normal mapping**
  - Phong shader we are assuming that the surface of the polygon is smoothly curved.
  - One solution would be to increase the number of polygons to represent all the deformities, but this is computationally unfeasible for most applications.
  - Instead we use textures called normal maps to simulate minor perturbations in the surface normal.
  - Rather than arrays of colours, normal maps can be considered as arrays of **vectors**. These vectors are added to the interpolated normals to give the appearance of roughness.
  - **Pro**
    - Provide the illusion of surface texture
  - **Cons**
    - Does not affect silhouette
    - Does not affect occlusion calculation

## Rasterisation

- Rasterisation is the process of converting lines and polygons represented by their vertices into fragments.
  - Fragments are like pixels but include color, depth, texture coordinate. They may also never make it to the screen due to hidden surface removal or culling.
  - This operation needs to be accurate and efficient.
- 

- **Bresenham's algorithm**

- The key idea is that calculations are done incrementally, based on the values for the previous pixel.
- assume to begin with that the line is in the first octant.
- For each x we work out which pixel we set next
  - The next pixel with the same y value if the line passes below the midpoint between

the two pixels


- Or the next pixel with an increased y value if the line passes above the midpoint between the two pixels

```
int y = y0;
for (int x = x0; x <= x1; x++) {
    setPixel(x,y);
    M = (x + 1, y + 1/2)
    if (M is below the line)
        y++
}
```

```
int y = y0;
int w = x1 - x0;
int h = y1 - y0;
int F = 2 * h - w;

for (int x = x0; x <= x1; x++) {
    drawPixel(x,y);
    if (F < 0)
        F += 2*h;
    else {
        F += 2*(h-w); y++;
    }
}
```

○ Example



$$w = 8$$

$$h = 5$$

$$2 * (h - w) = -6$$

$$2 * h = 10$$

| x | y | F  |
|---|---|----|
| 0 | 0 | 2  |
| 1 | 1 | -4 |
| 2 | 1 | 6  |
| 3 | 2 | 0  |
| 4 | 3 | -6 |
| 5 | 3 | 4  |
| 6 | 4 | -2 |
| 7 | 4 | 8  |
| 8 | 5 | 2  |

● Polygon filling


○ Shared edges

- We adopt a rule: The edge pixels belong to the rightmost and/ or upper polygon;

i.e. do not draw rightmost or uppermost edge pixels


- **Scanline algorithm**

- Testing every pixel is very inefficient.
- We only need to check where the result changes value, i.e. when we cross an edge
- We proceed row by row:
  - Calculate intersections incrementally.
  - Sort by x value.
  - Fill runs of pixels between intersections.
- **Edge table**
  - The edge table is a lookup table indexed on the y-value of the lower vertex of the edge.
  - Horizontal edges are not added
  - We store the the x-value of the lower vertex, the increment (inverse gradient) of the edge and the y-value of the upper vertex.



- **Active Edge List**

- On the left edge, round up to the nearest integer, with  $\text{round}(n) = n$  if  $n$  is an integer.
- On the right edge, round down to the nearest integer, but with  $\text{round}(n) = n-1$  if  $n$  is an integer.



- **Antialiasing**

- **Prefiltering**

- Prefiltering is computing exact pixel values geometrically rather than by sampling.

| Input  | Result |     |     |     |     |     |     |     |     |     |     |     |     |     |     |   |   |   |   |  |  |  |  |  |  |  |  |  |     |  |  |  |  |  |  |  |  |
|--|--------|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|---|---|---|---|--|--|--|--|--|--|--|--|--|-----|--|--|--|--|--|--|--|--|
| <table><tr><td>0</td><td>0</td><td>0</td><td>0.2</td><td>0.7</td><td>0.5</td></tr><tr><td>0.1</td><td>0.4</td><td>0.8</td><td>0.9</td><td>0.5</td><td>0.1</td></tr><tr><td>0.5</td><td>0.7</td><td>0.3</td><td>0</td><td>0</td><td>0</td></tr></table> | 0      | 0   | 0   | 0.2 | 0.7 | 0.5 | 0.1 | 0.4 | 0.8 | 0.9 | 0.5 | 0.1 | 0.5 | 0.7 | 0.3 | 0 | 0 | 0 | <table><tr><td></td><td></td><td></td><td></td><td></td><td></td></tr><tr><td></td><td></td><td></td><td>0.9</td><td></td><td></td></tr><tr><td></td><td></td><td></td><td></td><td></td><td></td></tr></table> |  |  |  |  |  |  |  |  |  | 0.9 |  |  |  |  |  |  |  |  |
| 0  | 0      | 0   | 0.2 | 0.7 | 0.5 |     |     |     |     |     |     |     |     |     |     |   |   |   |   |  |  |  |  |  |  |  |  |  |     |  |  |  |  |  |  |  |  |
| 0.1  | 0.4    | 0.8 | 0.9 | 0.5 | 0.1 |     |     |     |     |     |     |     |     |     |     |   |   |   |   |  |  |  |  |  |  |  |  |  |     |  |  |  |  |  |  |  |  |
| 0.5  | 0.7    | 0.3 | 0   | 0   | 0   |     |     |     |     |     |     |     |     |     |     |   |   |   |   |  |  |  |  |  |  |  |  |  |     |  |  |  |  |  |  |  |  |
|  |        |     |     |     |     |     |     |     |     |     |     |     |     |     |     |   |   |   |   |  |  |  |  |  |  |  |  |  |     |  |  |  |  |  |  |  |  |
|  |        |     | 0.9 |     |     |     |     |     |     |     |     |     |     |     |     |   |   |   |   |  |  |  |  |  |  |  |  |  |     |  |  |  |  |  |  |  |  |
|  |        |     |     |     |     |     |     |     |     |     |     |     |     |     |     |   |   |   |   |  |  |  |  |  |  |  |  |  |     |  |  |  |  |  |  |  |  |

- **Prefiltering is most accurate but requires more computation.**

- **Postfiltering**

- **Postfiltering can be faster. Accuracy depends on how many samples are taken per pixel. More samples means larger memory usage.**
    - Postfiltering is taking samples at a higher resolution (supersampling) and then averaging.
    - Draw the line at a higher resolution and average (supersampling).


| Higher resolution | Sample | Output |
|-------------------|--------|--------|
|                   |        |        |

- **Weighted postfiltering**

- It is common to apply weights to the samples to favour values in the center of the pixel.

- **Stochastic sampling** (Random)


- Taking supersamples in a grid still tends to produce noticeably regular aliasing effects.
- Adding small amounts of jitter to the sampled points makes aliasing effects appear as visual noise.
- **Adaptive Sampling**
  - Supersampling in large areas of uniform colour is wasteful.
  - Supersampling is most useful in areas of major colour change.
  - Solution: Sample recursively, at finer levels of detail in areas with more colour variance.



## Particle systems (粒子系统)

- Some visual phenomena are best modelled as collections of small particles.
- Particles are usually represented as small textured quads or point sprites – single vertices with an image attached.
- They are billboarded, i.e transformed so that they are always face towards the camera.
- Particles are created by an emitter object and evolve over time, usually changing position, size, colour.





## Global Lighting

---

- The lighting equation we looked at earlier only handled direct lighting from sources:
  - $I = I_a \rho_a + \sum_{l \in \text{lights}} I_l (\rho_d (\widehat{S}_1 \cdot \widehat{m}) + \rho_{sp} (\widehat{r}_1 \cdot \widehat{v})^f)$
  - We added an ambient fudge term to account for all other light in the scene.
  - Without this term, surfaces not facing a light source are black.
  - Methods that take this kind of multi-bounce lighting into account are called global lighting methods.

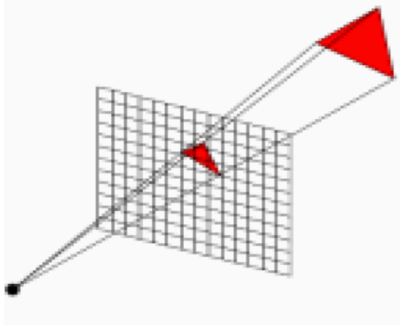
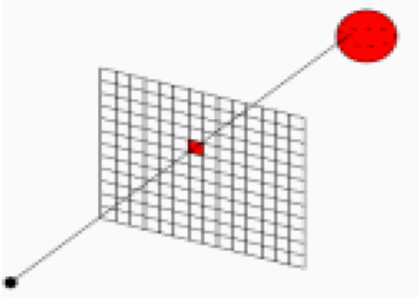
## Raytracing (Global lighting)

---

- Raytracing models specular reflection and refraction.
- Both methods are computationally expensive and are rarely suitable for real-time rendering.
- Ray tracing is a different approach to rendering than the pipeline we have seen so far.
- **Projective Methods vs RayTracing**
  - **Projective Method**
    - For each **object**: Find and update each pixel it influences
  - **Ray Tracing**
    - For each **pixel**: Find each object that influences it and update accordingly
  - Same
    - shading models
    - calculation of intersections,
  - Difference


- projection and hidden surface removal come for 'free' in ray tracing

○

| Projective  | Ray Tracing  |
|---|--|
|  |  |
|   | Same   |

- Location of pixel


○



- $pixelWidth = \frac{2w}{c}$
- $pixelHeight = \frac{2h}{r}$
- $i_c = -w + x(\frac{2w}{c}) = w(\frac{2x}{c} - 1)$
- $j_r = h(\frac{2y}{r} - 1)$
- The point  $P(x,y)$  of pixel  $(x,y)$  is given by:
  - $P(x,y) = E + w(\frac{2x}{c} - 1)i + h(\frac{2y}{r} - 1)k - nk$
- A ray from the camera through  $P(x,y)$  is given by:
  - $R(t) = E + t(P(x,y) - E) = E + tv$
  - $v = w(\frac{2x}{c} - 1)i + h(\frac{2y}{r} - 1)k - nk$

- $t = 0$ , we get E (Eye/Camera)
- $t = 1$ , we get  $P(x,y)$  – the point on the near plane
- $t > 1$  point in the world
- $t < 0$  point behind the camera – not on ray

○ **Generic Sphere**




$$t_{1,2} = \frac{-B \pm \sqrt{B^2 - 4AC}}{2A}$$

○ **Shadows**


- At each hit point we cast a new ray towards each light source. These rays are called shadow feelers.

○ **Extents (大小)**


- Extents are bounding boxes or spheres which enclose an object
- To compute a box extent for a mesh we simply take the min and max x, y and z coordinates over all the points.
- To compute a sphere extent we find the centroid of all the vertices by averaging their coordinates. This is the centre of the sphere.




Good fit



Better fit



Good fit




Poor fit

- **Projection extents**

- A projection extent of an object is a bounding box which encloses all the pixels which would be in the image of the object (ignoring occlusions).

- **Binary Space Partitioning (BSP)**

- Another approach to optimisation is to build a Binary Space Partitioning (BSP) tree dividing the world into cells, where each cell contains a small number of objects.



- **Raytracing Can't Do**

- Basic recursive raytracing cannot do:
    - Light bouncing off a shiny surface like a mirror and illuminating a diffuse surface

- Light bouncing off one diffuse surface to illuminate others
  - Light transmitting then diffusing internally
- Also a problem for rough specular reflection
  - Fuzzy reflections in rough shiny objects
- **Realtime ray-tracing (RTX)**
  - Works by arranging objects in a bounding volume hierarchy (BVH)
  - Specialised hardware offers fast traversal of these hierarchies to find ray intersections.

## Radiosity (Global lighting)

---

- Radiosity models diffuse reflection
- Radiosity is a global illumination technique which performs indirect diffuse lighting.
- Direct lighting techniques only take into account light coming directly from a source.
- Raytracing takes into account specular reflections of other objects.
- Radiosity takes into account diffuse reflections of everything else in the scene.
- **Finite elements**
  - We divide the scene up into small patches.
  - We then calculate the energy transfer from each patch to every other patch.
  - **Energy transfer**
    - The basic equation for energy transfer is:
    - Light output = Light emitted +  $\rho$  \* Light input
    - Where  $\rho$  is the diffuse reflection coefficient.
    - $B_i = E_i + \rho_i \sum_j B_j F_{ij}$ 
      - $B_i$  is the radiosity of patch i
      - $E_i$  is the energy emitted by patch i
      - $\rho_i$  is the reflectivity of patch i
      - $F_{ij}$  is a form factor which encodes what fraction of light from patch j reaches patch i.

## Color

---

- **RGB**
  - Red, Green, Blue
  - Colour is expressed as the addition of red, green and blue components.
- **CMYK**
  - Cyan, Magenta, Yellow, Black
  - CMY is a subtractive colour model, typically used in describing printed media.
- **HSV**

- Hue, Saturation, Value (Brightness)
- HSV (aka HSB) is an attempt to describe colours in terms that have more perceptual meaning
- H represents the hue as an angle from 0° (red) to 360° (red)
- S represents the saturation from 0 (grey) to 1 (full colour)
- V represents the value/brightness from 0 (black) to 1 (bright colour).

- **HSL**

- Hue, Saturation, Lightness
- HSL (aka HLS) replaces the brightness parameter with a (perhaps) more intuitive lightness value.
- H represents the hue as an angle from 0° (red) to 360° (red)
- S represents the saturation from 0 (grey) to 1 (full colour)
- L represents the lightness from 0 (black) to 1 (white).