

COMP342 I

Textures 2 and Rasterisation

Robert Clifton-Everest

Email: robertce@cse.unsw.edu.au

lostandtaken.com

opengameart.org/textures

textures.com

Rendering to a texture

A common trick is to set up a camera in a scene, render the scene into an offscreen buffer, then copy the image into a texture to use as part of another scene.


E.g. Implementing a security camera in a game.

//In OpenGL you can use

```
gl.glCopyTexImage2D (...);
```

Reflection


- To do better quality reflections we need to compute where the reflected light is coming from in the scene.



Reflection mapping


- Doing this in general is expensive, but we can make a reasonable approximation with textures:
 - Generate a cube that encloses the reflective object.
 - Place a camera at the centre of the cube and render the outside world onto the faces of the cube.
 - Use this image to texture the object

Reflection mapping




Reflective
object

Reflection mapping




Cube


Reflection mapping




Reflection mapping



Reflection mapping




Reflection mapping



Reflection mapping

- To apply the reflection-mapped texture to the object we need to calculate appropriate texture coordinates.
- We do this by tracing a ray from the camera, reflecting it off the object and then calculating where it intersects the cube.

Reflection mapping



Reflection mapping

- **Pros:**


- Produces reasonably convincing polished metal surfaces and mirrors

Reflection mapping

- Cons:
 - Expensive: Requires 6 additional render passes per object
 - Angles to near objects are wrong.
 - Does not handle self-reflections or recursive reflections.

OpenGL

- Cube maps can also be used for sky boxes.




OpenGL

- OpenGL has built in support for fast approximate reflection mapping (cube mapping).
- See ModelViewer
- OpenGL also has sphere mapping support, although this usually produces more distortion and is not as effective as cube mapping.


Shadows

- Our lighting model does not currently produce shadows.
- We need to take into account whether the light source is occluded by another object.



Shadow buffering

- One solution is to keep a **shadow buffer** for each light source.
- The shadow buffer is like the depth buffer, it records the distance from the light source to the closest object in each direction.



Shadow buffer

- Shadow rendering is usually done in multiple passes:
 1. Render the scene from each light's viewpoint capturing only z-info in shadow (depth) buffer (color buffer turned off)
 2. Render the scene from camera's point of view, using the previously captured shadow buffers to modulate the fragments

Shadow buffer

- When rendering a point P :
 - Project the point into the light's clip space.
 - Calculate the index (i,j) for P in the shadow buffer
 - Calculate the pseudodepth d relative to the light source
 - If $\text{shadow}[i,j] < d$ then P is in the shadow

Shadow buffer

- **Pros:**
 - Provides realistic shadows
 - No knowledge or processing of the scene geometry is required

Shadow buffer

- **Cons:**

- More computation
- Shadow quality is limited by precision of shadow buffer. This may cause some aliasing artefacts.
- Shadow edges are hard.
- The scene geometry must be rendered once per light in order to generate the shadow map for a spotlight, and more times for an omnidirectional point light.



OpenGL

- <http://www.paulsprojects.net/tutorials/smt/smt.html>

Light Mapping

- If our light sources and large portions of the geometry are static then we can precompute the lighting equations and store the results in textures called **light maps**.
- This process is known as **baked lighting**.

Light Mapping

- **Pros:**

- Sophisticated lighting effects can be computed at compile time, where speed is less of an issue.

Light mapping

- **Cons:**

- Memory and loading times for many individual light maps.
- Not suitable for dynamic lights or moving objects.
- Potential aliasing effects depending on the resolution of the light maps.

Normal mapping


- When we interpolate normals in a Phong shader we are assuming that the surface of the polygon is smoothly curved.
- What if the surface is actually rough with many small deformities?
- Putting a rough texture on a smooth flat surface looks wrong.

Normal mapping

- One solution would be to increase the number of polygons to represent all the deformities, but this is computationally unfeasible for most applications.
- Instead we use textures called **normal maps** to simulate minor perturbations in the surface normal.


Normal maps

- Rather than arrays of colours, normal maps can be considered as arrays of **vectors**. These vectors are added to the interpolated normals to give the appearance of roughness.




Normal maps

- Rather than arrays of colours, normal maps can be considered as arrays of **vectors**. These vectors are added to the interpolated normals to give the appearance of roughness.



Normal maps


- Rather than arrays of colours, normal maps can be considered as arrays of **vectors**. These vectors are added to the interpolated normals to give the appearance of roughness.



Normal map

Normal maps

- Rather than arrays of colours, normal maps can be considered as arrays of **vectors**. These vectors are added to the interpolated normals to give the appearance of roughness.





Perturbed
normals

Normal maps

- Pros:
 - Provide the illusion of surface texture
- Cons:
 - Does not affect silhouette
 - Does not affect occlusion calculation

Normal Mapping



OpenGL

- <http://www.opengl-tutorial.org/intermediate-tutorials/tutorial-13-normal-mapping/>
- <https://hacksoflife.blogspot.com/2009/11/per-pixel-tangent-space-normal-mapping.html>
- <http://www.terathon.com/code/tangent.html>


Rasterisation

- **Rasterisation** is the process of converting lines and polygons represented by their **vertices** into **fragments**.
- **Fragments** are like pixels but include color, depth, texture coordinate. They may also never make it to the screen due to hidden surface removal or culling.

Rasterisation

- This operation needs to be **accurate** and **efficient**.
- For this reason we prefer to use simple integer calculations.
- All our calculations are now in 2D screen space.

Drawing lines




Drawing lines - **bad**

```
float m = (y1-y0) / (float) (x1-x0) ;  
float b = y0 - m * x0 ;  
  
for (int x = x0; x <= x1; x++) {  
    int y = (int) Math.round(m * x + b) ;  
    drawPixel(x, y) ;  
}
```

Problems

- Floating point math is slow and creates rounding errors
 - Floating point multiplication, addition and round for each pixel
- Code does not consider:
 - Points are not connected if $m > 1$
 - Divide by zero if $x_0 == x_1$ (vertical lines)
 - Doesn't work if $x_0 > x_1$

Example: $y = 2x$



Incremental – still bad

```
// incremental algorithm
float m = (y1-y0) / (float) (x1-x0);
float y = y0;


for (int x = x0; x <= x1; x++) {
    y += m; //one less multiplication
    drawPixel(x, (int) Math.round(y));
}
```

Bresenham's algorithm

- We want to draw lines using only **integer calculations** and avoid multiplications.
- Such an algorithm is suitable for fast implementation in hardware.
- The key idea is that calculations are done **incrementally**, based on the values for the previous pixel.

Bresenham's algorithm


- We shall assume to begin with that the line is in the **first octant**.
- I.e. $x_1 > x_0$, $y_1 > y_0$ and $m \leq 1$



Bresenham's Idea

- For each x we work out which pixel we set next
 - The next pixel with the same y value if the line passes below the midpoint between the two pixels
 - Or the next pixel with an increased y value if the line passes above the midpoint between the two pixels

Bresenham's algorithm



Pseudocode

```
int y = y0;
for (int x = x0; x <= x1; x++) {
    setPixel(x,y);
    M = (x + 1, y + 1/2)
    if (M is below the line)
        y++
}
```

Testing above/below

- We're on the line when:

$$m = \frac{y - y_0}{x - x_0}$$

$$y - y_0 = m(x - x_0)$$

$$0 = m(x - x_0) - (y - y_0)$$

Testing above/below

- We're below the line when:

$$0 < m(x - x_0) - (y - y_0)$$

$$0 < (h / w)(x - x_0) - (y - y_0)$$

$$0 < h(x - x_0) - w(y - y_0)$$

$$0 < 2h(x - x_0) - 2w(y - y_0)$$

Testing above/below


We call this value F

$$F(x, y) = 2h(x - x_0) - 2w(y - y_0)$$

$$F(x, y) < 0 \Rightarrow (x, y) \text{ is above line}$$

$$F(x, y) > 0 \Rightarrow (x, y) \text{ is below line}$$

Midpoints



Incrementally

$$\begin{aligned} F(M) &= 2h(x_0 + 1 - x_0) - 2w(y_0 + \tfrac{1}{2} - y_0) \\ &= 2h - w \end{aligned}$$

$$\begin{aligned} F(M_1) &= 2h(x_0 + 2 - x_0) - 2w(y_0 + \tfrac{1}{2} - y_0) \\ &= F(M) + 2h \end{aligned}$$


$$\begin{aligned} F(M_2) &= 2h(x_0 + 2 - x_0) - 2w(y_0 + \tfrac{3}{2} - y_0) \\ &= F(M) + 2h - 2w \end{aligned}$$

Complete

```
int y = y0;
int w = x1 - x0; int h = y1 - y0;
int F = 2 * h - w;

for (int x = x0; x <= x1; x++) {
    drawPixel(x,y);
    if (F < 0) F += 2*h;
    else {
        F += 2*(h-w); y++;
    }
}
```

Example




$$w = 8$$
$$h = 5$$

```
int F = 2 * h - w;
```

x	y	F
0	0	2


Example



$$w = 8$$
$$h = 5$$

x	y	F
0	0	2

Example




$$w = 8$$

$$h = 5$$

$$2 * (h - w) = -6$$

x	y	F
0	0	2
1	1	-4

Example



$$W = 8$$


h = 5

$$2 * (h - w) = -6$$

$$2 * h = 10$$

x	y	F
0	0	2
1	1	-4
2	1	6

Example



$$w = 8$$


$$h = 5$$

$$2 * (h - w) = -6$$

$$2 * h = 10$$

x	y	F
0	0	2
1	1	-4
2	1	6
3	2	0

Example



$$w = 8$$


$$h = 5$$

$$2 * (h - w) = -6$$

$$2 * h = 10$$

x	y	F
0	0	2
1	1	-4
2	1	6
3	2	0
4	3	-6

Example



$$w = 8$$

$$h = 5$$

$$2 * (h - w) = -6$$

$$2 * h = 10$$


x	y	F
0	0	2
1	1	-4
2	1	6
3	2	0
4	3	-6
5	3	4
6	4	-2
7	4	8
8	5	2

Relaxing restrictions


- Lines in the other quadrants can be drawn by symmetrical versions of the algorithm.
- We need to be careful that drawing from P to Q and from Q to P set the same pixels.
- Horizontal and vertical lines are common enough to warrant their own optimised code.

Polygon filling

- Determining which pixels are inside a polygon is a matter of applying the **edge-crossing test** (from week 3) for each possible pixel.



Point in polygon




Problems

- Shared edges
 - If the pixel is on the edge of a polygon, should we draw it?
 - What if two polygons are adjacent (e.g. if they form part of a mesh)?
- Performance
 - Do we have to do the edge crossing test for every pixel?

Shared edges

- Pixels on shared edges between polygons need to be drawn consistently regardless of the order the polygons are drawn, with no gaps.
- We adopt a rule:
 - The edge pixels belong to the rightmost and/or upper polygon; i.e. do not draw rightmost or uppermost edge pixels




Scanline algorithm

- Testing every pixel is very inefficient.
- We only need to check where the result **changes value**, i.e. when we cross an edge.
- We proceed row by row:
 - Calculate intersections incrementally.
 - Sort by x value.
 - Fill runs of pixels between intersections.

Edge table

- The **edge table** is a lookup table indexed on the y-value of the lower vertex of the edge.
- Horizontal edges are not added
- We store the the x-value of the lower vertex, the increment (inverse gradient) of the edge and the y-value of the upper vertex.

Example



Edge table

y in	x	inc	y out
0	1	-0.25	4
0	5	1	1
0	9	-3	1
0	9	-0.4	5
3	2	-2	4
3	2	2.5	5

Active Edge List

- We keep a list of active edges that overlap the current scanline.
- Edges are added to the list as we pass the bottom vertex.
- Edges are removed from the list as we pass the top vertex.
- The edge intersection is updated incrementally.

Edges

- For each edge in the AEL we store:
 - The x value of its crossing with the current row (initially the bottom x value)
 - The increment (inverse gradient)
 - The y value of the top vertex.


```
//For every scanline
for (y = minY; y <= maxY; y++){


    remove all edges that end at y

    for (Edge e : active) {
        e.x = e.x + e.inc;
    }

    add all edges that start at y - keep list
sorted by x

    for (int i=0; i < active.size; i+=2){
        fillPixels(active[i].x, active[i+1].x,y);
    }
}
```


Example



Active edge list

x	inc	y out
1	-0.25	4
5	1	1
9	-3	1
9	-0.4	5

Example




Active edge list

x	inc	y out
1	-0.25	4
5	1	1
9	-3	1
9	-0.4	5

Example

y=1




Active edge list

x	inc	y out
0.75	-0.25	4
8.6	-0.4	5

Example


y=1



Active edge list

x	inc	y out
0.75	-0.25	4
8.6	-0.4	5


Example



Active edge list

x	inc	y out
0.5	-0.25	4
8.2	-0.4	5

Example




Active edge list

x	inc	y out
0.5	-0.25	4
8.2	-0.4	5

Example

y=3




Active edge list

x	inc	y out
0.25	-0.25	4
2	-2	4
2	2.5	5
7.8	-0.4	5

Example

y=3




Active edge list

x	inc	y out
0.25	-0.25	4
2	-2	4
2	2.5	5
7.8	-0.4	5

Example

y=4




Active edge list

x	inc	y out
4.5	2.5	5
7.4	-0.4	5

Example

y=4




Active edge list

x	inc	y out
4.5	2.5	5
7.4	-0.4	5

Example

y=5



Active edge list


x	inc	y out

OpenGL

- OpenGL is optimised for implementation on hardware.
- Hardware implementations do not work well with variable length lists.
- If polygons are **convex** the active edge list always has 2 entries.


Aliasing

- Lines and polygons drawn with these algorithms tend to look **jagged** if the pixel size is too large.
- This is another form of **aliasing**.



Aliasing

- Lines and polygons drawn with these algorithms tend to look **jagged** if the pixel size is too large.
- This is another form of **aliasing**.

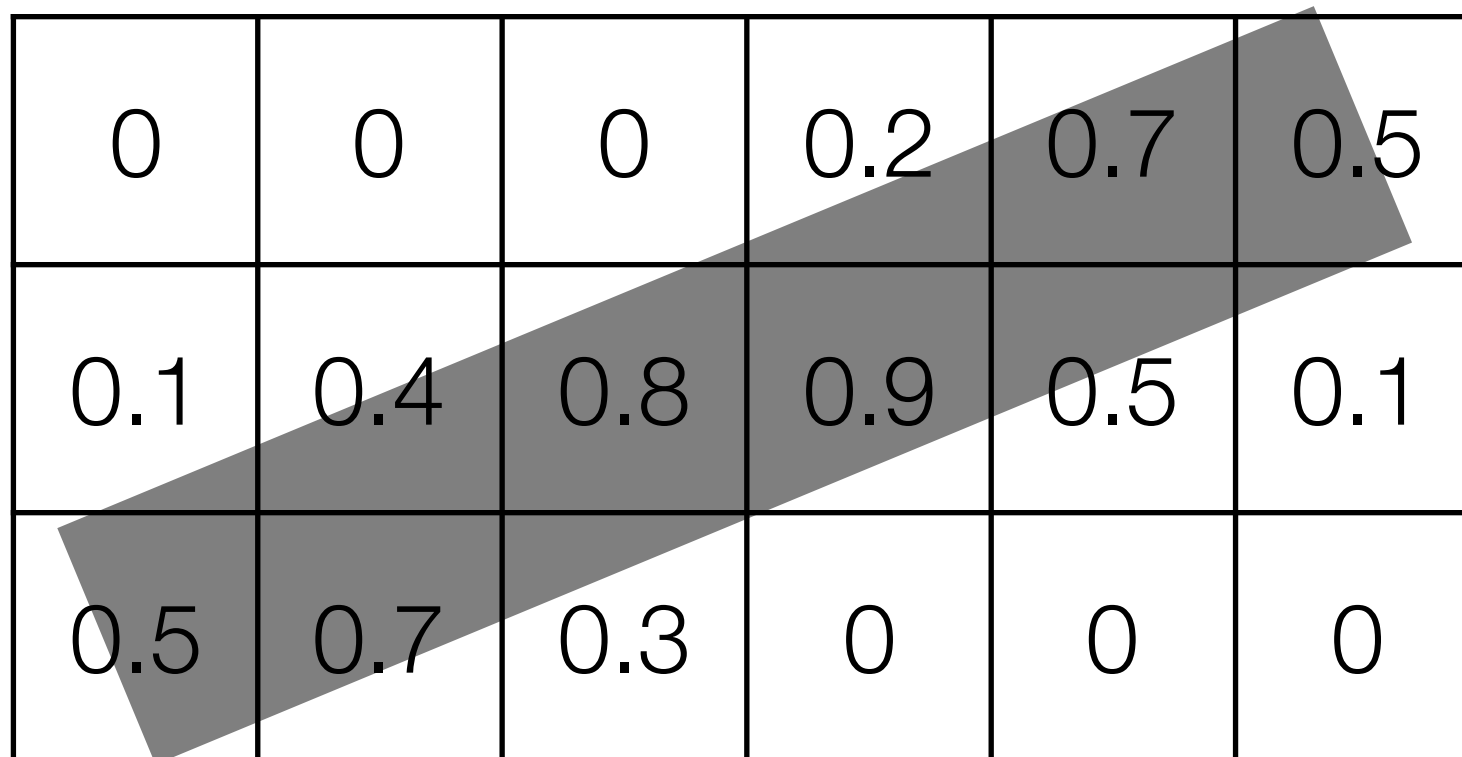


Antialiasing

- There are two basic approaches to eliminating aliasing (**antialiasing**).
- **Prefiltering** is computing exact pixel values geometrically rather than by sampling.
- **Postfiltering** is taking samples at a higher resolution (supersampling) and then averaging.

Prefiltering

- For each pixel, compute the amount occupied and set pixel value to that percentage.




A 3x6 grid of pixel values. A gray diagonal band highlights the cells where the value is greater than 0.5. The values in the grid are:

0	0	0	0.2	0.7	0.5
0.1	0.4	0.8	0.9	0.5	0.1
0.5	0.7	0.3	0	0	0


Prefiltering

- For each pixel, compute the amount occupied and set pixel value to that percentage.




Postfiltering

- Draw the line at a higher resolution and average (supersampling).




Postfiltering

- Draw the line at a higher resolution and average (supersampling)




Postfiltering

- Draw the line at a higher resolution and average (supersampling)



Comparing

Prefiltering



Postfiltering




Weighted postfiltering

- It is common to apply weights to the samples to favour values in the center of the pixel.

$1/16$	$1/16$	$1/16$
$1/16$	$1/2$	$1/16$
$1/16$	$1/16$	$1/16$

Stochastic sampling


- Taking supersamples in a grid still tends to produce noticeably regular aliasing effects.
- Adding small amounts of **jitter** to the sampled points makes aliasing effects appear as visual noise.




Adaptive Sampling

- Supersampling in large areas of uniform colour is wasteful.
- Supersampling is most useful in areas of major colour change.
- Solution: Sample recursively, at finer levels of detail in areas with more colour variance.


Adaptive sampling



Adaptive sampling



Adaptive sampling



Antialiasing

- Prefiltering is most accurate but requires more computation.
- Postfiltering can be faster. Accuracy depends on how many samples are taken per pixel. More samples means larger memory usage.

OpenGL

// implementation dependant, may not
even do anything 😊

```
gl.glEnable(GL3.GL_LINE_SMOOTH);  
gl.glHint(GL3.GL_LINE_SMOOTH_HINT, GL3.GL  
_NICEST);
```

// also requires alpha blending

```
gl.glEnable(GL2.GL_BLEND);  
gl.glBlendFunc(GL2.GL_SRC_ALPHA,  
               GL2.GL_ONE_MINUS_SRC_ALPHA);
```

OpenGL

```
// full-screen multi-sampling
GLCapabilities capabilities =
    new GLCapabilities();
capabilities.setNumSamples(4);
capabilities.setSampleBuffers(true);
...
gl.glEnable(GL.GL_MULTISAMPLE);
```