

# COMP342 I

---

## Ray Tracing

Robert Clifton-Everest

Email: [robertce@cse.unsw.edu.au](mailto:robertce@cse.unsw.edu.au)

# Recap: Global Lighting

---

- The lighting equation we looked at earlier only handled **direct lighting** from sources:

$$I = \boxed{I_a \rho_a} + \sum_{l \in \text{lights}} I_l \left( \rho_d (\hat{\mathbf{s}}_l \cdot \hat{\mathbf{m}}) + \rho_{sp} (\hat{\mathbf{r}}_l \cdot \hat{\mathbf{v}})^f \right)$$

- We added an **ambient fudge term** to account for all other light in the scene.
- Without this term, surfaces not facing a light source are black.

# Recap: Global lighting

---

- In reality, the light falling on a surface comes from **everywhere**. Light from one surface is reflected onto another surface and then another, and another, and...
- Methods that take this kind of multi-bounce lighting into account are called **global lighting** methods.


# Raytracing and Radiosity

---

- There are two main methods for global lighting:
  - **Raytracing** models specular reflection and refraction.
  - **Radiosity** models diffuse reflection.
- Both methods are **computationally expensive** and are rarely suitable for real-time rendering.

# Rays

---



# Shadows

---

- We can add shadows very simply.
- At each hit point we cast a new ray towards each light source. These rays are called **shadow feelers**.
- If a shadow feeler intersects an object before it reaches the source, then omit that source from the illumination equation for the point.


# Self-shadows

---

- We need to take care when the shadow is cast by the hit object itself.
- The shadow feeler will always intersect the hit object at time  $t=0$ .
- This intersection is only relevant if the light is on the opposite side of the object.

# Example


---






# Example

---




# Example

---




# Example

---



# Example

---



# PseudoCode


---

```
Trace primary ray
if (hit is null)
    set (x,y) to background
else
    set (x,y) = ambient color
    Trace secondary ray to each light
        if not blocked from light
            (x,y) += contribution from light
```

# Reflections

---

- We can now implement realistic reflections by casting further reflected rays.




$$\mathbf{r} = \mathbf{v} - 2(\mathbf{v} \cdot \mathbf{n})\mathbf{n}$$

# Reflections

---


- Reflected rays can in turn be reflected off another object and another.
- We usually write our code to stop after a fixed number of reflections to avoid infinite recursion.



# Transparency

---

- We can also model transparent objects by casting a second ray that continues through the object.






# Transparency

---

- Transparency can also be applied reflexively, yielding a tree of rays.



# Illumination

---

- The illumination equation is extended to include reflected and transmitted components, which are computed recursively:


$$I(P) = I_{amb} + I_{dif} + I_{spe} + I(P_{ref}) + I(P_{tra})$$

- We will need material coefficients to attenuate the reflected and transmitted components appropriately.

# Refraction of Light

---

- When a light ray strikes a transparent object, a portion of the ray penetrates the object. The ray will change direction from **dir** to **t** if the speed of light is different in medium 1 and medium 2. Vector **t** lies in the same plane as **dir** and the normal **m**.



# Refraction

---

- To handle transparency appropriately we need to take into account the refraction of light.
- Light bends as it moves from one medium to another. The change is described by Snell's Law:

$$\frac{\sin \theta_1}{c_1} = \frac{\sin \theta_2}{c_2}$$

- where  $c_1$  and  $c_2$  are the speeds of light in each medium.

# Example Snell's law


---

Air:

$$c_1 = 99.97\% c$$

Glass:

$$c_2 \approx 55\% c$$



# Example


---

- Suppose medium 2 is some form of glass in which light only travels 55% as fast as in medium 1 which is the air. Suppose the angle of incidence of the light is 60 degrees from the normal. What is the angle of the transmitted light?


# Refraction

- The figure (a) shows light moving from the faster medium to the slower, and (b) shows light moving from the slower to the faster medium.
- The angles pair together in the same way in both cases; only the names change.


a)




b)



c)




d)



# Refraction

- In (c) and (d), the larger angle has become nearly  $90^\circ$ . The smaller angle is near the **critical angle**: when the smaller angle (of the slower medium) gets large enough, it forces the larger angle to  $90^\circ$ . A larger value is impossible, so no light is transmitted into the second medium. This is called **total internal reflection**.






# Refraction

---

- Different wavelengths of light move at different speeds (except in a vacuum).
- So for maximum realism, we should calculate different paths for different colours.



# Refraction of Light

---

- Simplest to model transparent objects so that their index of refraction does not depend on wavelength.
- To do otherwise would require tracing separate rays for each of the color components, as they would refract in somewhat different directions.
- This would be expensive computationally, and would still provide only an approximation, because an accurate model of refraction should take into account a large number of colors, not just the three primaries.

# Optimisation

---

- Testing collisions for more complex shapes (such as meshes) can be very time consuming.
- In a large scene, most rays will not hit the object, so performing multiple expensive collision tests is wasteful.
- We want fast ways to rule out objects which will not be hit.


# Extents

---


- Extents are bounding boxes or spheres which enclose an object.
- Testing against a box or sphere is fast.
- If this test succeeds, then we proceed to test against the object.
- We want tight fitting extents to minimise false positives.

# Extents


---




Good fit



Better fit



Good fit



Poor fit

# Computing extents

---

- To compute a **box extent** for a mesh we simply take the min and max x, y and z coordinates over all the points.
- To compute a **sphere extent** we find the **centroid** of all the vertices by averaging their coordinates. This is the centre of the sphere.
- The radius is the distance to the vertex farthest from this point.

# Projection extents


---

- Alternatively, we can build extents in **screen space** rather than **world space**.
- A **projection extent** of an object is a bounding box which encloses all the pixels which would be in the image of the object (ignoring occlusions).
- Pixels outside this box can ignore the object.
- Doesn't work for shadow feelers or reflected rays

# Projection extents

---

- We can compute a projection extent of a mesh by projecting all the vertices into screen space and finding the min and max x and y values.






# Binary Space Partitioning

---


- Another approach to optimisation is to build a Binary Space Partitioning (BSP) tree dividing the world into cells, where each cell contains a small number of objects.



# BSPs

---


- Another approach to optimisation is to build a BSP tree dividing the world into cells, where each cell contains a small number of objects.



# BSPs

---


- Another approach to optimisation is to build a BSP tree dividing the world into cells, where each cell contains a small number of objects.



# BSPs

---


- Another approach to optimisation is to build a BSP tree dividing the world into cells, where each cell contains a small number of objects.



# BSPs

---


- Another approach to optimisation is to build a BSP tree dividing the world into cells, where each cell contains a small number of objects.



# BSPs

---


- Another approach to optimisation is to build a BSP tree dividing the world into cells, where each cell contains a small number of objects.



# BSPs


---

- Another approach to optimisation is to build a BSP tree dividing the world into cells, where each cell contains a small number of objects.



# BSPs


- Another approach to optimisation is to build a BSP tree dividing the world into cells, where each cell contains a small number of objects.





# Traversing the tree

- In this case we do not want to traverse the entire tree. We only want to visit the leaves the ray passes through.



# Traversal algorithm



---

```
visit(E, v, node): (E eye)
  if (node is leaf):
    intersect ray with objs in leaf
  else:
    if (E on left):
      visit(E, v, left)
      other = right;
    else:
      visit(E, v, right)
      other = left
    endif



    if (ray crosses boundary):
      E' = intersect(E, v, boundary)
      visit(E', v, other)
    endif

  endif
```



# Traversing the tree





# Traversing the tree





# Traversing the tree





# Traversing the tree



# Traversing the tree





# Traversing the tree









# Traversing the tree





# Traversing the tree





# Traversing the tree





# Traversing the tree





# Traversing the tree





# Traversing the tree



# Traversing the tree





# Traversing the tree







# Traversing the tree



# Traversing the tree



# Exercise

---

- How does milk look different to white paint?

# Exercise


---

- How does milk look different to white paint?
- Both are opaque.
- Both are essentially pure white.
- Milk is an example of scattering

# Scattering

---


- Scattering (or subsurface scattering) is when light refracts into an object that is non-uniform in its density and is reflected out at a different angle and position.



# Scattering

---

- Milk is a substance that has this property.
- As is skin, leaves, and wax.
- Typically, they are hard to render.



# Scattering

---



# Scattering

---

- We don't really have time to cover this in more depth in this course. Read this if you want to know more (NOT EXAMINABLE).
- <http://graphics.ucsd.edu/~henrik/images/subsurf.html>



# Raytracing Can't Do

---

- Basic recursive raytracing cannot do:
  - Light bouncing off a shiny surface like a mirror and illuminating a diffuse surface
  - Light bouncing off one diffuse surface to illuminate others
  - Light transmitting then diffusing internally
- Also a problem for rough specular reflection
  - Fuzzy reflections in rough shiny objects

# Realtime ray-tracing (RTX)

---

- Recent Nvidia innovation
- Realtime raytracing supported via specialised hardware
- Programs have to be written to use it, but it fits in quite well with existing graphics pipelines

# Realtime ray-tracing (RTX)

---

- Works by arranging objects in a **bounding volume hierarchy** (BVH)
- Specialised hardware offers fast traversal of these hierarchies to find ray intersections.

# Bounding Volume Hierarchy (BVH)

- Scene is divided into small volumes.
- Child volumes are contained entirely within their parents.
- Sibling volumes may overlap







Figure 2.2

Letters a through g represent the objects.  
Numbers 1 through 5 represent interior nodes.

# Bounding Volume Hierarchy (BVH)

- Parents can have an arbitrary number of children
- The best way to divide up the scene depends on a lot of factors



*Figure 2.2*

Letters a through g represent the objects.  
Numbers 1 through 5 represent interior nodes.

# Top down construction

---

- Divide the set of primitives in the scene into two (or more) subsets then recursively subdivide those until all subsets contain only one primitive.
- Easy to implement.
- Usually faster than alternatives.
- Doesn't always produce the best possible tree.

# Bottom up construction

---

- Treat all primitives in the scene as a set of leaves. Pick two (or more) of them and group them into a node. Repeat till there is only one node in the set.
- Slower than top-down in most cases.
- More difficult to implement.
- Tends to produce better trees.

# RTX Examples

---

- <https://www.youtube.com/watch?v=WoQr0k2IA9A>
- <https://www.youtube.com/watch?v=Ms7d-3Dprio>
- [https://www.youtube.com/watch?v=1IliQZw\\_p\\_E](https://www.youtube.com/watch?v=1IliQZw_p_E)



# Volumetric ray tracing

---

- We can also apply ray tracing to **volumetric** objects like smoke or fog or fire.
- Such objects are transparent, but have different intensity and transparency throughout the volume.

# Volumetric Ray Tracing

---

- We represent the volume as two functions:

$C(P)$  = colour at point  $P$


$\alpha(P)$  = transparency at point  $P$

- Typically these are represented as values in a 3D array. Interpolation is used to find values at intermediate points.
- These functions may in turn be computed based on density, lighting or other physical properties.

# Sampling

---

- We cast a ray from the camera through the volume and take samples at fixed intervals along the ray.



# Sampling

---

- We end up with  $(N+1)$  samples:

$$P_i = R(t_{hit} + i\Delta t)$$

$$C_i = C(P_i)$$

$$\alpha_i = \alpha(P_i)$$

$$C_N = (r, g, b)_{background}$$

$$\alpha_N = 1$$

# Alpha compositing

---

- We now combine these values into a single colour by applying the alpha-blending equation.

$$C_N^N = C_N$$

$$C_N^i = \alpha_i C_i + (1 - \alpha_i) C_N^{i+1}$$

Total  
colour  
at i

Local  
colour  
at i

Total  
colour  
at i+1

# Exercise

---

- Suppose we have a background color of  $(0, 1, 0)$  and a volume with the uniform color of  $(1, 0.5, 0.5)$ . A ray cast through that volume takes two samples. The first has an alpha value of 0.2 and the second 0.1. What is the colour of the resulting pixel?

# Alpha compositing

---

- We can write a closed formula for the colour from a to b as:

$$C_b^a = \sum_{i=a}^b \alpha_i C_i \prod_{j=a}^{i-1} (1 - \alpha_j)$$

- We can compute this function from front to back, stopping early if the transparency term gets small enough that nothing more can be seen.

# In OpenGL

---

- Volumetric ray tracing (aka ray casting) does not require a full ray tracing engine.
- It can be implemented in OpenGL as a fragment shader applied to a cube with a 3D texture.
- <https://www.shadertoy.com/view/XsIGRr>



# Examples

---

- <https://www.shadertoy.com/view/Ms2SD1>
- <https://www.shadertoy.com/view/4sS3zG>
- See <http://shadertoy.com/> for more examples.

# Sources

---

- [http://en.wikipedia.org/wiki/Volume ray casting](http://en.wikipedia.org/wiki/Volume_ray_casting)
- [http://graphics.ethz.ch/teaching/former/scivis\\_07/Notes/Slides/03-raycasting.pdf](http://graphics.ethz.ch/teaching/former/scivis_07/Notes/Slides/03-raycasting.pdf)
- [http://http.developer.nvidia.com/GPUGems/gpugems\\_ch39.html](http://http.developer.nvidia.com/GPUGems/gpugems_ch39.html)